

---

# Super Resolution On Global Weather Forecasts

---

Bryan Zhang, Dhruv Rao, Adam Yang, Lawrence Zhang, Rodz Andrie Amor  
University of Maryland  
Department of Computer Science

## 1 Introduction

Weather forecasting is a vitally important tool for tasks ranging from planning day to day activities to disaster response planning. However, modeling weather has proven to be a challenging task due to its chaotic and unpredictable nature. Each variable, from temperature to precipitation to wind, all influence the path the environment will take. As a result, all models tend to rapidly lose accuracy as the temporal range of their forecasts increase.

Classical forecasting methods use a myriad of physics-based, numerical, and stochastic techniques to predict the change in weather variables over time. However, such forecasts often require a very large amount of data and are extremely computationally expensive. Furthermore, as climate and global weather patterns change, classical models are substantially more difficult and time-consuming to update for changing environments.

Fortunately, with recent advances in deep learning and publicly available high quality weather datasets, deploying learning methods for estimating these complex systems has become feasible. The current state-of-the-art deep learning models have comparable accuracy to the industry standard numerical models and are becoming more ubiquitous in practice due to their adaptability.

Recently, GraphCast [3] demonstrated the ability to produce fast and accurate global weather predictions for up to 10 days into the future, using only very recent weather data as input. This has proven to be a monumental breakthrough for weather prediction research, especially with weather neural networks. However, while GraphCast has proven its usefulness for global weather forecasts, it does not have as much utility on a smaller scale because of its low resolution. This limits its value for the kind of day-to-day localized weather forecasting that the vast majority of people would find useful.

Our group seeks to improve upon existing deep learning based forecasting methods by increasing spatial resolutions of global weather predictions. Specifically, we are interested in performing super resolution (SR) on GraphCast temperature predictions by increasing the global precision from 1 degree of accuracy to 0.5 degrees, which is approximately 111km and 55km respectively.

## 2 Motivation

Global forecasts provide high level trends of climate throughout the entire planet. While such forecasts can model long range trends in weather, the results are too coarse for local weather applications. For instance, a model that produces weather forecast images at 1 degree spatial resolution will have pixels that cover over 100km. While much work has been done to create global forecasting models that have increased spatial and temporal accuracy, there has not been a high volume of work dedicated to modifying existing models to better meet the demands of localized forecasts. In this paper, we explore using deep learning methods to increase the spatial resolution of coarse global forecasts by increasing the precision of image forecasts produced by Graphcast

This task can be described as performing SR and downscaling from low resolution to high resolution. Inspired by a paper on deep convolutional SR networks [1], we propose using a similar SR model that is specifically trained to artificially downscale forecasts at higher spatial resolution. Instead of training a larger forecasting model that come at the cost of energy and training time, lightweight SR

39 models as a cheap alternative to extract weather features to create higher resolution predictions. This  
40 offers the benefit of the fast global weather forecasting with the higher granularity of larger models.

### 41 **3 Prior Work and State of the Art**

#### 42 **3.1 GraphCast [3]**

43 GraphCast is the current standard for graph-based neural network global weather forecasting model.  
44 It is a common benchmark for weather forecasts because of its relatively low computation requirement  
45 and ability to produce a global medium-range 10-day forecast in under a minute on a Google Cloud  
46 Tensor Processing Unit. The model uses an autoregressive approach to make further date forecasts,  
47 utilizing its own previous six-hour forecast as the input by projecting variables onto a mesh and  
48 decoding learned features back onto a standard latitude-longitude grid. To evaluate the accuracy of  
49 the predictions, the researchers calculated the root mean squared error and the anomaly correlation  
50 coefficient to that of other models with the HRES dataset as the baseline. GraphCast demonstrated a  
51 statistically significant improvement in accuracy in all weather variables, and it subsequently became  
52 the benchmark standard machine learning weather prediction model. Using a frozen GraphCast, we  
53 are forwarding weather data to create forecasts on temperature, precipitation, and cloud coverage  
54 which will then be upscaled.

#### 55 **3.2 Image Super-Resolution Using Deep Convolutional Networks [1]**

56 This paper proposes a computationally lightweight but remarkably effective convolutional neural  
57 network that transforms a low resolution image to a high resolution image. The SRCNN consists  
58 of three major steps: patch extraction, non-linear mapping, and reconstruction. It learns the end-  
59 to-end mapping from low resolution inputs to high resolution outputs. This is especially useful for  
60 applications in deblurring or upscaling images. Finally, the SRCNN architecture performs well under  
61 various benchmarks as well as metrics like MSE and Peak Signal-to-Noise Ratio (PSNR). We use  
62 this architecture to perform the image SR from the global weather forecast data to a higher detail.  
63 However, this paper uses a 3 channel mapping to represent RGB whereas global weather data comes  
64 in 1 channel, which can represent a variable like temperature or cloud coverage. We have modified  
65 the architecture used in this paper to work with one channel input and output, as well as added a  
66 series of residual blocks and skip connections to improve its performance.

#### 67 **3.3 Implicit Neural Representations with Periodic Activation Functions [7]**

68 This paper proposes a new approach to generating Implicit Neural Representations (INR) by using  
69 periodic activation functions. INR approximate a continuous function from data that is represented  
70 discretely through using simple fully connected Multi-Layer Perceptrons as universal approximation  
71 functions. The goal is for the Multi-Layer Perceptron to learn a mapping between a discrete input  
72 to a continuous output. Notably, this paper highlights the significant effectiveness of Multi-Layer  
73 Perceptrons as INRs when using a sine activation function (SIRENs) which is infinitely differentiable,  
74 with each derivative being a smooth function. Essentially, this means that SIRENs are good at  
75 modeling smooth surfaces and transitions. These INRs have direct applications within the field of SR  
76 due to their ability to map pixel coordinates directly to approximated values.

#### 77 **3.4 GAN based Super Resolution [4]**

78 This paper proposes a super-resolution generative adversarial network (SRGAN) that utilizes a deep  
79 residual network with skip connections. Rather than using MSE loss, the researchers devised a  
80 perceptual loss centered around high-level feature maps and a discriminator that tries to ensure  
81 generated images are similar to input ones. The perceptual loss is a weighted linear combination  
82 of standard content loss from most other deep learning networks and adversarial loss, or loss  
83 calculated from probabilities of the binary classification the discriminator is performing. The  
84 generator architecture consists of residual blocks followed by 2 convolutional layers that extract  
85 features and produce higher-resolution images, and the discriminator is a series of convolutional  
86 layers ending with a sigmoid function to extract binary classification probabilities. For testing, in  
87 addition to analyzing the content loss over time, the researchers employed human subjects to score

88 the quality of generated images from different models on a scale from 1 to 5. The researchers found  
89 that the model performs favorably on image datasets in comparison with other leading models of the  
90 time, such as ResNets.

## 91 4 Data

92 The base GraphCast forecasts numerous variables for a myriad atmospheric levels which are included  
93 in ERA5 data. We decided to select a few variables based on their relative importance for local areas.  
94 In particular, we selected, cloud coverage, and temperature 2 meters above sea level, as we believe  
95 that these features have the most impact on day-to-day life. Furthermore, each variable has its own  
96 challenges and advantages. We found that global cloud coverage had the best visualization and was  
97 the easiest to interpret visually, while temperature had the best qualities for testing upscaling due to  
98 its global patterns and relative spatio-temporal variance.

99 The GraphCast model is trained on ERA5 model, which makes it particularly suited to work with  
100 that dataset. ERA5 is incredibly detailed, well documented, and readily available for any use, which  
101 enabled us to experiment with the dataset [2]. Fortunately, with the structure the ERA5 dataset takes  
102 and the information GraphCast can make inputs on, we were able to create a one-to-one mapping  
103 from GraphCast output to the corresponding ERA5 ground truth.

104 In our data pipeline, the first step we implemented collects a large number of inferences from the base  
105 GraphCast output. By doing so, we do not have to perform a full GraphCast inference every time we  
106 want to load input data into a data loader. We limited inferences to only output our relevant features  
107 in order to reduce the size of the dataset. Next, we made a function that uses the date labeled on the  
108 output file to match it with the corresponding ERA5 file. Finally, we implemented a data loader that  
109 maps the GraphCast output with the exact corresponding ERA5 data to use as the input and ground  
110 truth pair.

## 111 5 Methods

112 Instead of modifying the architecture of GraphCast to work with different resolutions, we decided to  
113 maintain a frozen GraphCast and use it as a black box for the input of our model. Then, we developed  
114 a SR model that works with the output of base GraphCast. This enables us to easily modify the  
115 structure of our network without having to deal with the larger, complex GraphCast architecture. In  
116 order to obtain a dataset to train our models, we produced inferences from GraphCast over 12 hour  
117 intervals for 200 days and 2 images a day, totaling to 400 images. These inferences were saved in a  
118 Google Drive folder to be later pulled and used in a data loader for training, validating, and testing.  
119 Next, we use the time stamp label for each image to map to the corresponding ERA5 images as the  
120 ground truth. The objective of the model is to minimize the error when upscaling GraphCast image to  
121 ERA5 ground truth. For the loss functions, we experimented with mean squared error, mean absolute  
122 error, edge error, and perception error. In the architecture fine-tuning part of the process, our group  
123 has experimented with a variety of models and loss functions. We experimented on UNet, SRCNN,  
124 GAN, and INR, with varying levels of success. The following subsections will go into further detail  
125 on the process and evaluation of the architectures we implemented..

### 126 5.1 UNet

127 The UNet architecture [5] was previously developed for object segmentation tasks. The original paper  
128 that introduced it displayed use cases in biomedical image processing, but it was quickly adopted  
129 for a wide variety of other tasks. The high level architecture uses convolutional layers and pooling  
130 to produce a low level feature map of an image with decreased spatial dimension compared to the  
131 original image passed in. The model then up-samples the feature map input resolution back to the  
132 original spatial dimensions and uses a final softmax layer to produce segmentation results. For SR  
133 tasks, the output can be transformed to a regression objective where outputs model a continuous  
134 weather variable.

135 We decided to first experiment with the UNet architecture because of its relative simplicity compared  
136 to other deep learning models and our previous familiarity with it. We hoped to establish it as a  
137 relatively straightforward base before exploring other, more complex models. However, outputs of

the UNet model seemed to average the inputs and were unable to capture original image features. The output of the UNet simply was uninterpretable and showed no sign of capturing the patterns and upscaling the weather data. As a result, we started experimenting with other models listed below.

## 5.2 SRGAN

The next model we investigated was a SR generative adversarial network (SRGAN). The idea behind a GAN is to have two competing networks being trained at the same time, also known as adversarial training [4]. The generator network generates new images by making modifications to the input data, whereas the discriminator network tries to predict whether the generated data came from the input dataset or not. Ideally, the goal is to have the generator producing data that the discriminator cannot tell apart from the actual input. This has immediate benefits in image generation tasks, where we can assess the quality of generated images by seeing if the discriminator can distinguish them from input images, as well as from the accuracy of the prediction. In the context of SR models, the generator would be identical to the SRCNN architecture discussed in the next section, but now the discriminator can be used to verify the quality of the images the SRCNN creates. This discriminator utilizes convolutional layers to extract features from the generated image before performing a binary classification using the sigmoid activation function to classify the image as real input or generated, first proposed by Ledig, et al. (2016)[4]. If the discriminator cannot tell the difference between the two, the SRCNN generator component has produced a high-quality SR. However, we had recurring problems with implementing the GAN and ensuring that the model produced high quality and detailed images. In our case, the generator could not create any image complex enough to fool the discriminator, and the discriminator maintained a very low loss and high accuracy throughout the entire training process. With this in mind, we pivoted to experimenting with training the SRCNN with MSE error in the next section.

## 5.3 SRCNN

Finally, our best performing model is the SR convolutional neural network (SRCNN). This was a novel approach first devised by Dong, et al. (2014) [1] that utilizes a relatively simple architecture of 3 main layers to perform image SR. After a bicubic interpolation to upscale the input image to the desired size resulting in a pixelated low resolution map, the model extracts feature maps using convolutional layers, nonlinearly maps these feature maps to another set of high-dimensional feature blocks, and then constructs the final higher-resolution image. Later papers built on this baseline by making the model deeper using more convolutional layers, and then a paper by Tong, et al. (2017) [9] brought the idea of multiple skip connections into the fold to protect against vanishing gradients and overfitting. Work by Song, et al. (2019) [8] studied in more depth the idea of residual blocks, or networks where information could move from initial layers to end layers without any restriction through skip connections and then be combined with information that did pass through all the layers.

For our implementation, we found that the base SRCNN model from Dong, et al. worked remarkably well, but was not complex enough to capture the smaller details and specific patterns. As a result, in order to increase the model complexity, we have added residual blocks and skip connections. The intent of the skip connections is to preserve the features in later layers by directly bypassing the layers. Similarly, the residual blocks work almost identically to the skip connections, but the residuals blocks also contain a convolution that is intended to learn the small distinctions between the input and output. By adding these two layers into the base SRCNN, we were able to have it produce images in more detail and capture smaller patterns.

## 181 6 Results

182 To train our models, we used Mean Squared Error (MSE) loss to minimize the difference between  
 183 outputs and ground truth and as an evaluation metric. Across the three tested models, SRCNN  
 184 performed the best in matching the ground truth model according to our evaluation metric.




Figure 1: Results from UNet, SRGAN, and SRCNN Respectively

185 Additionally, we were able to demonstrate from ERA5 low resolution data of cloud coverage to a  
 186 higher resolution ERA5 cloud coverage. This demonstrates how the model is capable of downscaling  
 187 any type of weather data and is not restricted to global temperature.




Figure 2: Global Cloud Coverage Upscale




Figure 3: High Contrast Blowup of Image 80x80




Figure 4: High Contrast Blowup of Image 30x30

## 7 Super Resolution from INR with SIREN/WIRE

An alternative method for approaching SR tasks that doesn't utilize traditional CNNs is by using INRs to approximate a continuous function from data that is represented discretely. This has direct applications to our problem of upscaling the weather data from ERA5 as these continuous functions are able to make images of any resolution by interpolating on the function.

Our approach was to train multiple INRs using periodic and non-periodic activation functions. The use of periodic activation functions, such as sinusoidal and Gaussian, was explored by Sitzmann et al. [7] as an alternative to non-periodic activation functions, such as ReLU and Sigmoid, when designing an implicit representation.

### 7.1 Architecture




Figure 5: Non Linear Activation Functions

The underlying architecture of the INRs is the standard multi-layer perceptrons with 2 hidden layers, each hidden layer having 256 neurons. We trained 4 separate models using the same architecture but alternated between 4 different activations:


- ReLU (non-periodic)
- Gauss (periodic)
- Sinusoidal (periodic)
- Complex Gabor Wavelet (periodic)

### 7.2 Results


After training each INR model for 2000 epochs, the models using periodic activation functions (Gauss, Siren, Wire) significantly outperformed the model using the ReLU activation when trying to upscale an image of cloud formations from 64 x 128 to 256 x 512. These results are consistent with previous experiments by Sitzmann et al.[7]

When comparing across the different periodic activation functions using PSNR and SSIM as the evaluation metrics, the results are not consistent with Saragadam et al. [6]. Unlike their experiment, the SIREN model performed slightly better than WIRE here for our SR task, but this may be due to

the different attributes of weather images and the training images in the researcher’s dataset. When comparing the results of the INRs and SRCNN model, we found that the INRs had a lower MSE score across all the periodic activations.



(a) Results from INR Models



(b) SRCNN and Periodic INR results

The last experiment we ran with the INR was on 2m temperature (Temperature 2 meters above the surface level). We trained on the same SR task going from 64 x 128 to 256 x 512 and found that the WIRE and SIREN activations performed just as well as the cloud formation upscaling.

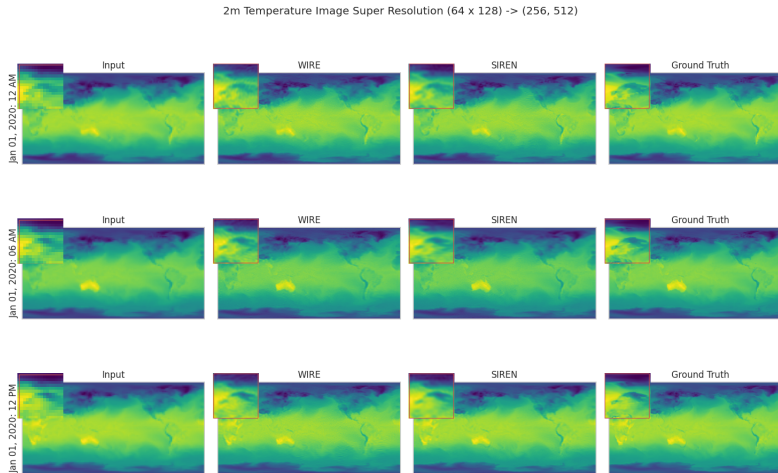


Figure 7: Surface Temperature SR

## 8 Future work

### 8.1 Multi Modal Data

While our models were able to produce reasonable estimations for a single variable, future work can seek SR methods that integrate multi-modal inputs. Image to image tasks are typical for SR objectives, but the high dimensionality of weather data offers a chance for more robust models trained to measure associations between various measurements. For our data pipeline, the architectures were only trained and tested on single channel images. However, a model capable of processing multimodal data will be able to recognize the correlations between variables and produce more fine-grain results.

## 227 8.2 Localized Models

228 While models such as GraphCast provide global forecasts, more lightweight and local models may  
229 be trained on limited areas. Using local models may prove more accurate for smaller regions while  
230 costing a fraction of running inference for global forecasts. Furthermore, local models may capture  
231 specific and nuanced patterns only seen by in an area that the global models are not suited to recognize.  
232 By restricting the boundary to a specific area and increasing the resolution, it will be substantially  
233 more informative for local residents. For very small areas, a precision of 0.1 degrees, or roughly 11  
234 kilometers, and even lower will certainly be significantly more useful.

## 235 8.3 Autoregressive SR

236 An additional method of SR is to approach the task in an auto-regressive style training scheme.  
237 Theoretically, this method would progressively denoise inputs and produce high fidelity outputs  
238 over a number of iterations. We demonstrate this with our SRCNN model, but found that training  
239 repeated iterations would be computationally expensive. In addition, the model is prone to introducing  
240 unwanted artifacts and remove accurate predictions over numerous iterations. However at a controlled  
241 amount, the model does potentially show increased precision and is especially useful in non-noisy  
242 regions especially




Figure 8: Autoregressive SRCNN



## References

- [1] Chao Dong, Chen Change Loy, Kaiming He, and Xiaoou Tang. Image super-resolution using deep convolutional networks. *IEEE transactions on pattern analysis and machine intelligence*, 38(2):295–307, 2015.
- [2] Hans Hersbach, Bill Bell, Paul Berrisford, Shoji Hirahara, András Horányi, Joaquín Muñoz-Sabater, Julien Nicolas, Carole Peubey, Raluca Radu, Dinand Schepers, Adrian Simmons, Cornel Soci, Saleh Abdalla, Xavier Abellan, Gianpaolo Balsamo, Peter Bechtold, Gionata Biavati, Jean Bidlot, Massimo Bonavita, Giovanna De Chiara, Per Dahlgren, Dick Dee, Michail Diamantakis, Rossana Dragani, Johannes Flemming, Richard Forbes, Manuel Fuentes, Alan Geer, Leo Haimberger, Sean Healy, Robin J. Hogan, Elías Hólm, Marta Janisková, Sarah Keeley, Patrick Laloyaux, Philippe Lopez, Cristina Lupu, Gabor Radnoti, Patricia de Rosnay, Iryna Rozum, Freja Vamborg, Sebastien Villaume, and Jean-Noël Thépaut. The era5 global reanalysis. *Quarterly Journal of the Royal Meteorological Society*, 146(730):1999–2049, 2020.
- [3] Remi Lam, Alvaro Sanchez-Gonzalez, Matthew Willson, Peter Wirnsberger, Meire Fortunato, Ferran Alet, Suman Ravuri, Timo Ewalds, Zach Eaton-Rosen, Weihua Hu, Alexander Merose, Stephan Hoyer, George Holland, Oriol Vinyals, Jacklynn Stott, Alexander Pritzel, Shakir Mohamed, and Peter Battaglia. Graphcast: Learning skillful medium-range global weather forecasting. 2023.
- [4] Christian Ledig, Lucas Theis, Ferenc Huszar, Jose Caballero, Andrew Cunningham, Alejandro Acosta, Andrew Aitken, Alykhan Tejani, Johannes Totz, Zehan Wang, and Wenzhe Shi. Photo-realistic single image super-resolution using a generative adversarial network. 2016.
- [5] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. 2015.
- [6] Vishwanath Saragadam, Daniel LeJeune, Jasper Tan, Guha Balakrishnan, Ashok Veeraraghavan, and Richard G. Baraniuk. Wire: Wavelet implicit neural representations, 2023.
- [7] Vincent Sitzmann, Julien N. P. Martel, Alexander W. Bergman, David B. Lindell, and Gordon Wetzstein. Implicit neural representations with periodic activation functions. 2020.
- [8] Dehua Song, Chang Xu, Xu Jia, Yiyi Chen, Chunjing Xu, and Yunhe Wang. Efficient residual dense block search for image super-resolution, 2019.
- [9] Tong Tong, Gen Li, Xiejie Liu, and Qinquan Gao. Image super-resolution using dense skip connections. pages 4809–4817, 2017.